



Pulsar

Realtime Analytics At Scale

Wang Xinglang



Agenda

- **Pulsar** : Real Time Analytics At eBay
 - Business Use Cases
 - Product Requirements
- **Pulsar** : Technology Deep Dive



Pulsar



Business Use Case: Behavioral Analytics

- Behavioral Analytics : Analysis Of User Behavior on eBay
 - Inputs : User Clicks, Impression, Conversions etc
 - Problem
 - In Session Targeting
 - Campaign Optimization

Requirements

- **Low Latency**
 - <1 Sec E2E latency
- **Scalability**
 - Millions of events/sec
- **Data Accuracy**
 - Data Sources are lossy



Business Use Case: Monitoring

- Monitoring : Application, Systems and Infrastructure Monitoring
 - Inputs : Heartbeats, SNMP, Logs
 - Problem
 - Collection, Aggregation of Metrics
 - Co-relation and Alerting

Requirements

- **Scalability**
 - 10s of Millions of events/sec
- **Availability**
 - 99.99% Uptime
- **Flexibility**
 - User Driven Rules
 - On the fly joins with other streams, data sources



Business Use Case: Security

- Security : Traffic Limiting, DOS Detection, Bot Detection
 - Inputs : Clicks, Impressions, Logs
 - Problem
 - Enforce Quotas in Real-time
 - Detect and Prevent Bad Bots

Requirements

- **Scalability**
 - 10s of Millions of events/sec
- **Latency**
 - 1 sec SLA
- **Flexibility**
 - User Drive Custom Rules
- **Data Accuracy**
 - 99% is acceptable



Pulsar Product Requirements Summary

- **Scalability**

- Scale to 10s of Millions Events/Sec

- **Latency**

- <1 Sec Delivery of Events

- **Availability**

- Highly Available System
- No downtime during upgrades
- Disaster Recovery Support across data centers

- **Flexibility**

- User Driven Complex Rules. Eg (CPU/TPS > 3 AND ERRORS/TPS>0.2)

- **Data Accuracy**

- Should deal with missing data
- 99.9% Delivery Guarantee

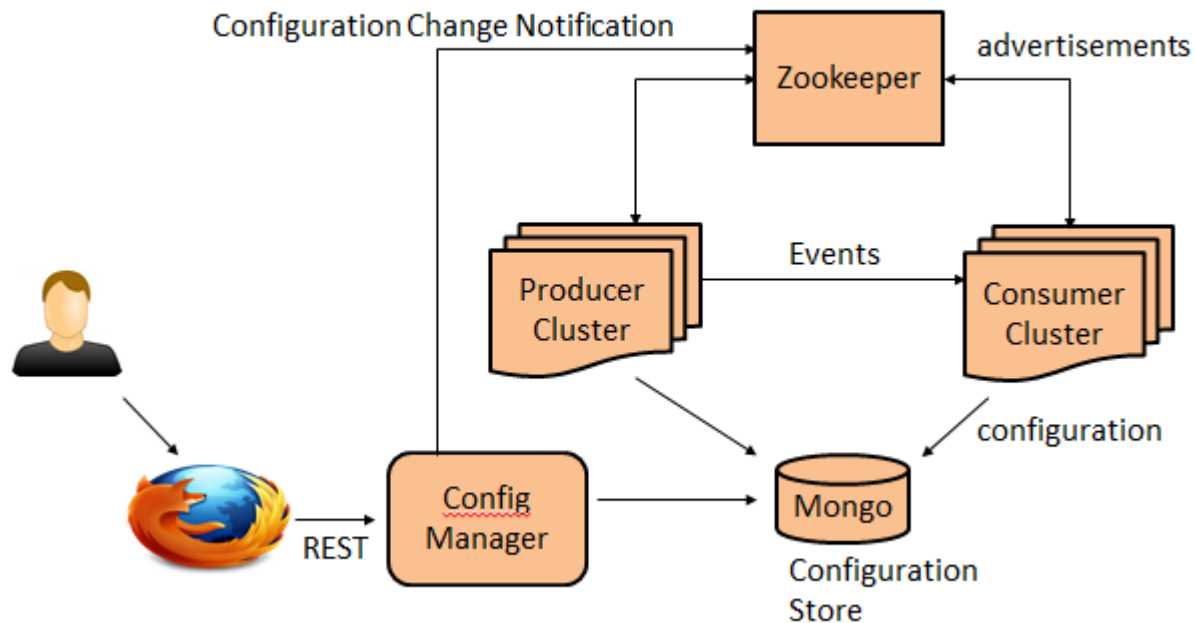


Pulsar Technical Deep Dive

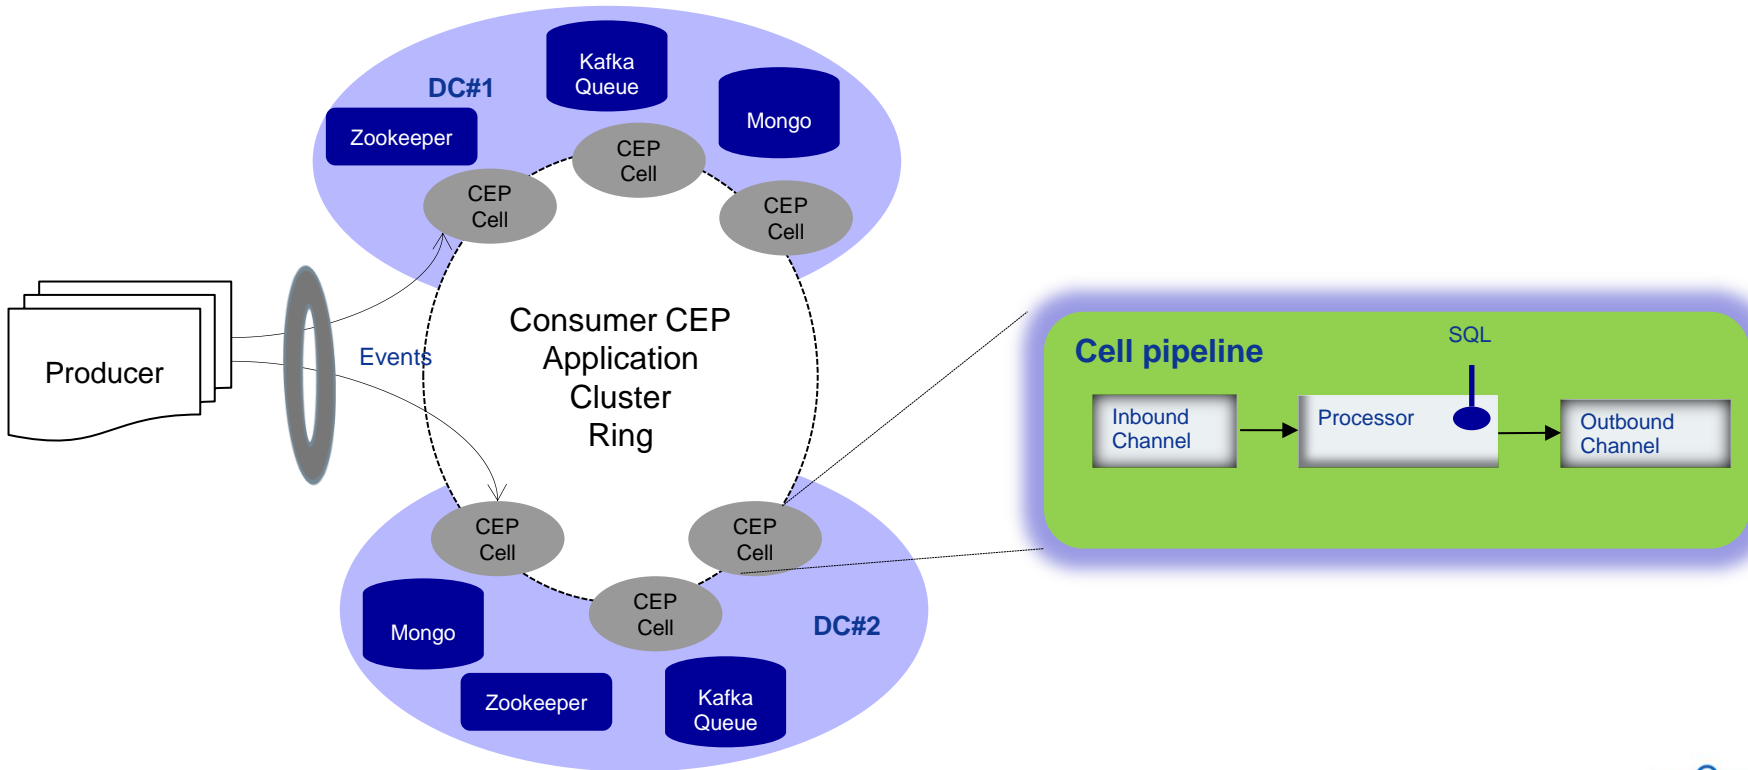


Pulsar CEP framework

JetStream

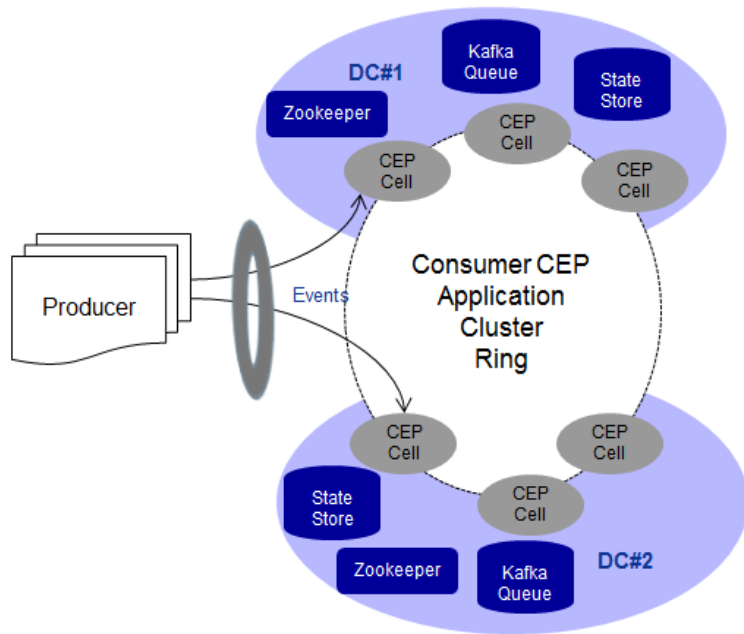


Pulsar Deployment Architecture





Availability And Scalability

- Multi datacenter failovers
- Dynamic Partitioning
- Elastic Clusters
- Self Healing
- Shutdown Orchestration
- Dynamic Flow Routing
- Dynamic Topology Changes



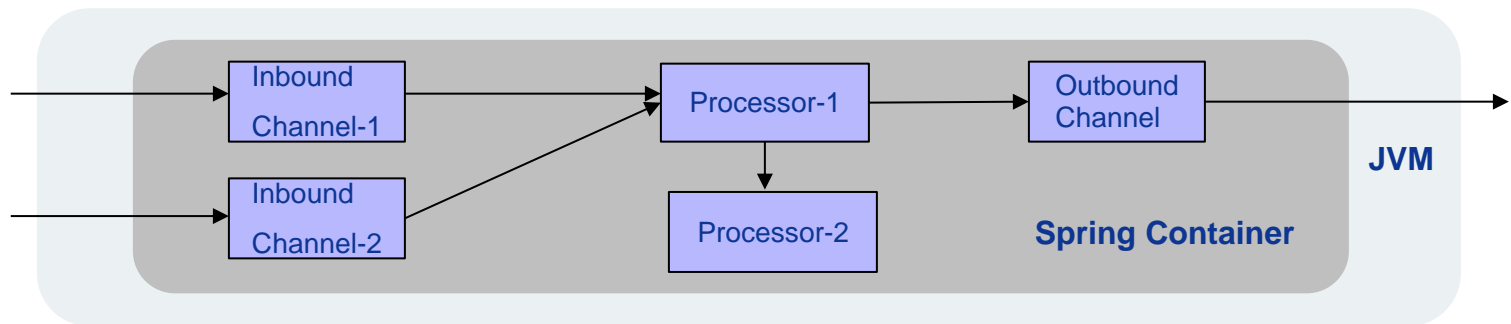
Pulsar Framework Building Blocks



- Event = Tuples (K,V) - Mutable
- Inbound Channel (Event Source) adapts to external world and sources events
- Outbound Channel (Event Sink) adapts to external world and consumes events
- Event Processor (Event Sink and Event Source)
- Pipelining, Batching & Flow control



Pulsar Application (CEP Cell)

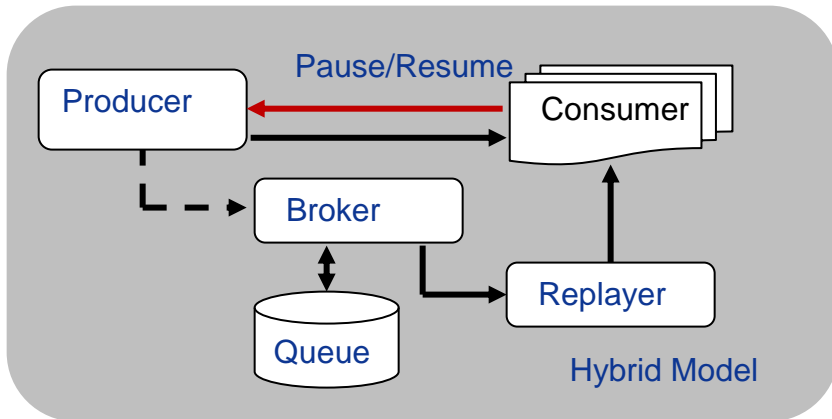

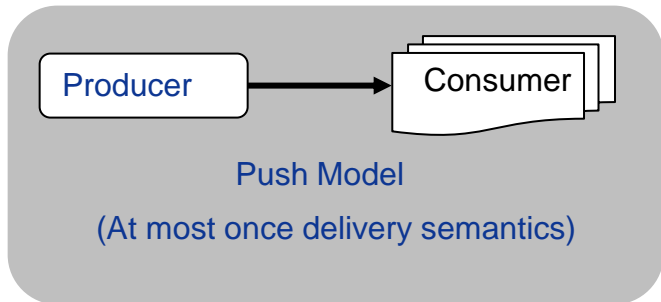


```
<bean id="InboundChannel-1"
class="com.ebay.jetstream.event.channel.messaging.InboundMessagingChannel">
    <property name="address" ref="InboundChannelAddress" />
    <property name="eventSinks">
        <list>
            <ref bean="Processor-1" />
        </list>
    </property>
</bean>
```

```
<bean id="Processor-1"
class="com.ebay.jetstream.event.processor.esper.EsperProcessor">
    <property name="esperEventListener" ref="EsperEventListener" />
    <property name="configuration" ref="EsperConfiguration" />
    <property name="epi" ref="EPL" />
    <property name="eventSinks">
        <list>
            <ref bean="outboundChannel" />
        </list>
    </property>
</bean>
```



Messaging And Clustering



Real Time Stream Processing

- The 8 Requirements of Real-Time Stream Processing (Ref. Stonebraker)
 - Rule 1: Keep the Data Moving
 - Rule 2: Query using SQL on Streams (StreamSQL)
 - Rule 3: Handle Stream Imperfections
 - Rule 4: Generate Predictable Outcomes
 - Rule 5: Integrate Stored and Streaming Data
 - Rule 6: Guarantee Data Safety and Availability
 - Rule 7: Partition and Scale Applications Automatically
 - Rule 8: Process and Respond Instantaneously
 - **Rule 9: Continuous Processing**
 - **Rule 10: Dynamic Topology Changes**
 - **Rule 11: Dynamic SQL changes & flow control**



Flexibility



Flexibility : CEP Engine (Esper)

- SQL like language for specifying processing rules
- Analysis over rolling and tumbling windows of time
- Filtering and Joining streams
- Grouping and Ordering output
- For routing events between stages and between clusters
- Event Mutation
- Correlation
- Patterns



Flexibility : DATA MODEL

```
<bean id="EventDefinition" class="com.ebay.jetstream.event.processor.esper.EsperDeclaredEvents">
  <property name="eventTypes">
    <list>
      <bean class="com.ebay.jetstream.event.processor.esper.MapEventType">
        <property name="eventAlias" value="RawEvent"/>
        <property name="eventFields">
          <map>
            <entry key="D1" value="java.lang.String"/>
            <entry key="D2" value="java.lang.String"/>
            <entry key="D3" value="java.lang.String"/>
            <entry key="D4" value="java.lang.Integer"/>
            <entry key="D5" value="java.lang.Integer"/>
            <entry key="createtime" value="java.lang.Long"/>
          </map>
        </property>
      </bean>
    </list>
  </property>
</bean>
```

key	value
js_ev_type	RawEvent
D1	"2045573"
D2	"79843743994"
D3	"908098404060"
D4	73840754
D5	1
createtime	7076303760



Flexibility : Event Processing Language (EPL)


Metrics over Rolling window

```
INSERT INTO RateOffender SELECT D1, count(*) AS total_count FROM RawEvent(D5 = 1).win:time(10 sec) group by D1 having count(*) > 8 output last every 1 second;
```

@OutputTo("OMC")

@PublishOn(topics="Rate.Monitor/rateOffender")

```
SELECT D1, 'block' as Policy FROM RateOffender;
```



Flexibility : Event Processing Language (EPL) – Filtering, Mutation & Routing


```
INSERT INTO STREAMROUTE SELECT D1, D2, D3, D4 FROM RawEvent(D1 in  
('2045573','2053742') and Common.isNumeric(D2) and D2 != '0');
```

```
@OutputTo("OMC")
```

```
@PublishOn(topics="ChanA/aEvent,ChanB/bEvent")
```

```
@ClusterAffinityTag(colname = "D2")
```

```
SELECT * FROM STREAMROUTE;
```



Flexibility : Event Processing Language (EPL) – Creating Multidimensional Metrics Over Tumbling Windows

```
create context MCContext start @now end after 10 seconds;
```

```
context MCContext
```

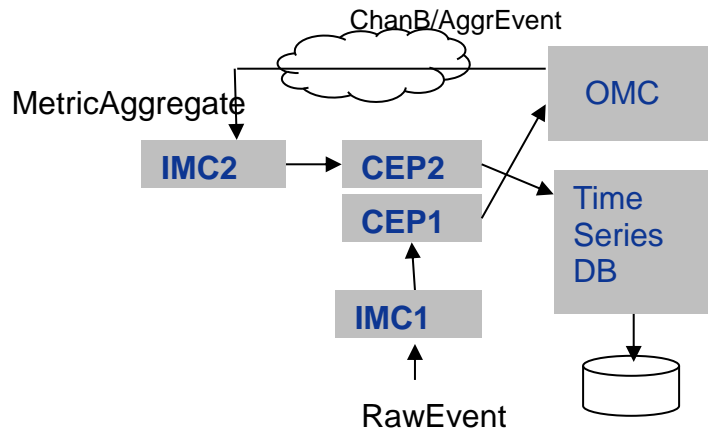
```
insert into MetricAggregate select count(*) as count, D1, D2, 'M1' as metricName from  
RawEvent(D1 is not null and D2 is not null) group by D1, D2 output snapshot when  
terminated;
```

```
@OutputTo("OMC")
```

```
@PublishOn(topics="ChanB/AggrEvent")
```

```
@ClusterAffinityTag(dimension=@CreateDimension(  
name="grpdim", dimensionspan="D1, D2, M1"))
```

```
select * from MetricAggregate;
```



Flexibility : Top N, Distinct Count & Percentiles Computation

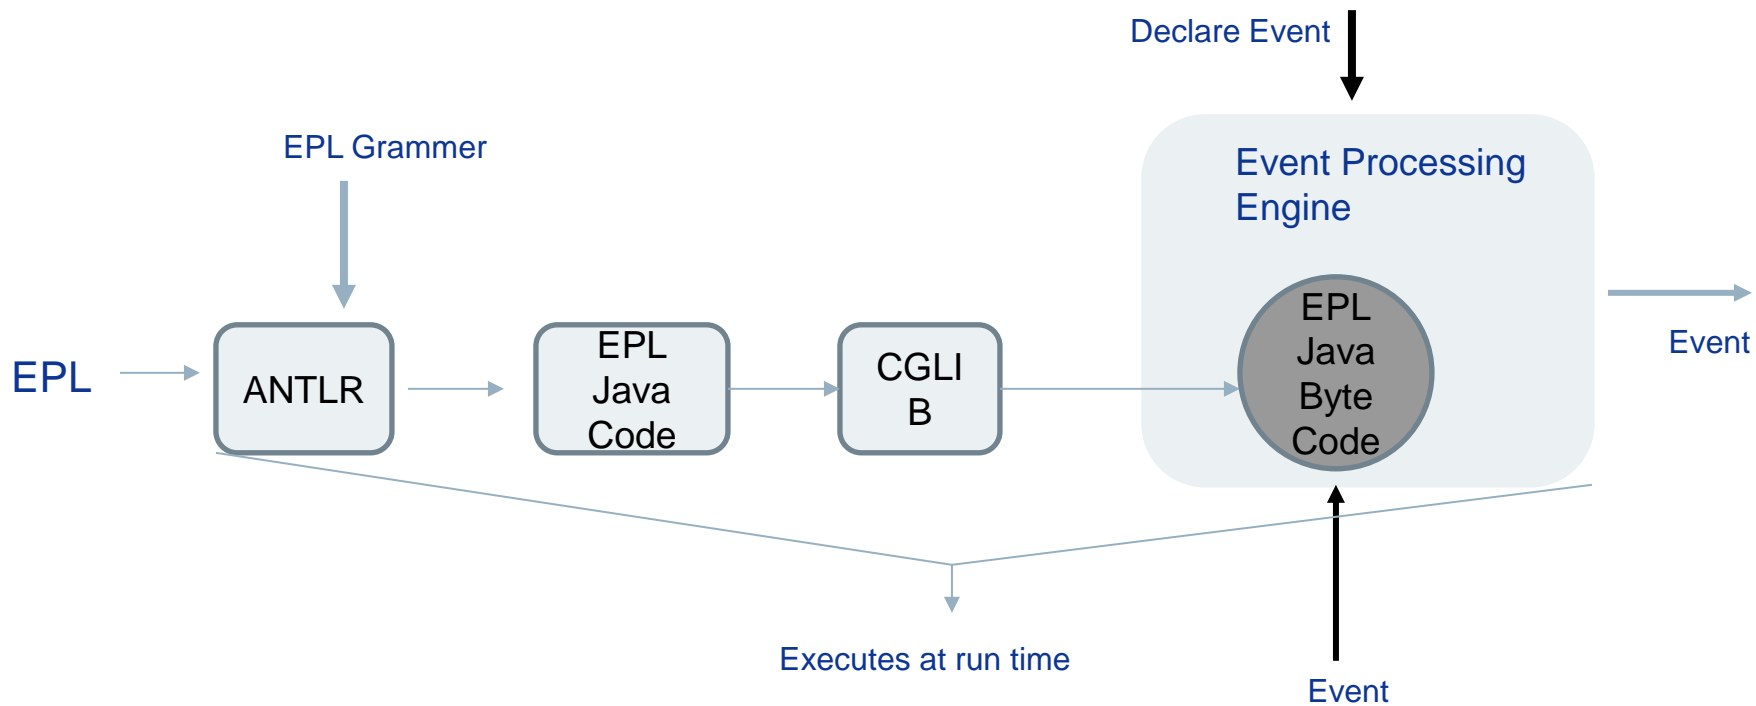
```
insert into TOPITEMSTREAM select count(*) as count, D1 from RawEvent() group by D1;  
select * from TOPITEMSTREAM order by count;
```

- Computations are heavy both in time and space.
- High Cardinality dimensions makes it worse
- Consider approximate algorithms
 - Margin of error – around 1%
 - Costs very little in space and time
 - Trade off accuracy for performance
- Implemented as aggregate functions

```
insert into TOPITEMSTREAM select TopN(1000, 10, D1) as topltems from RawEvent();  
select * from TOPITEMSTREAM;
```




Flexibility : Hot Deployment of EPL



Pulsar Stream Pipeline

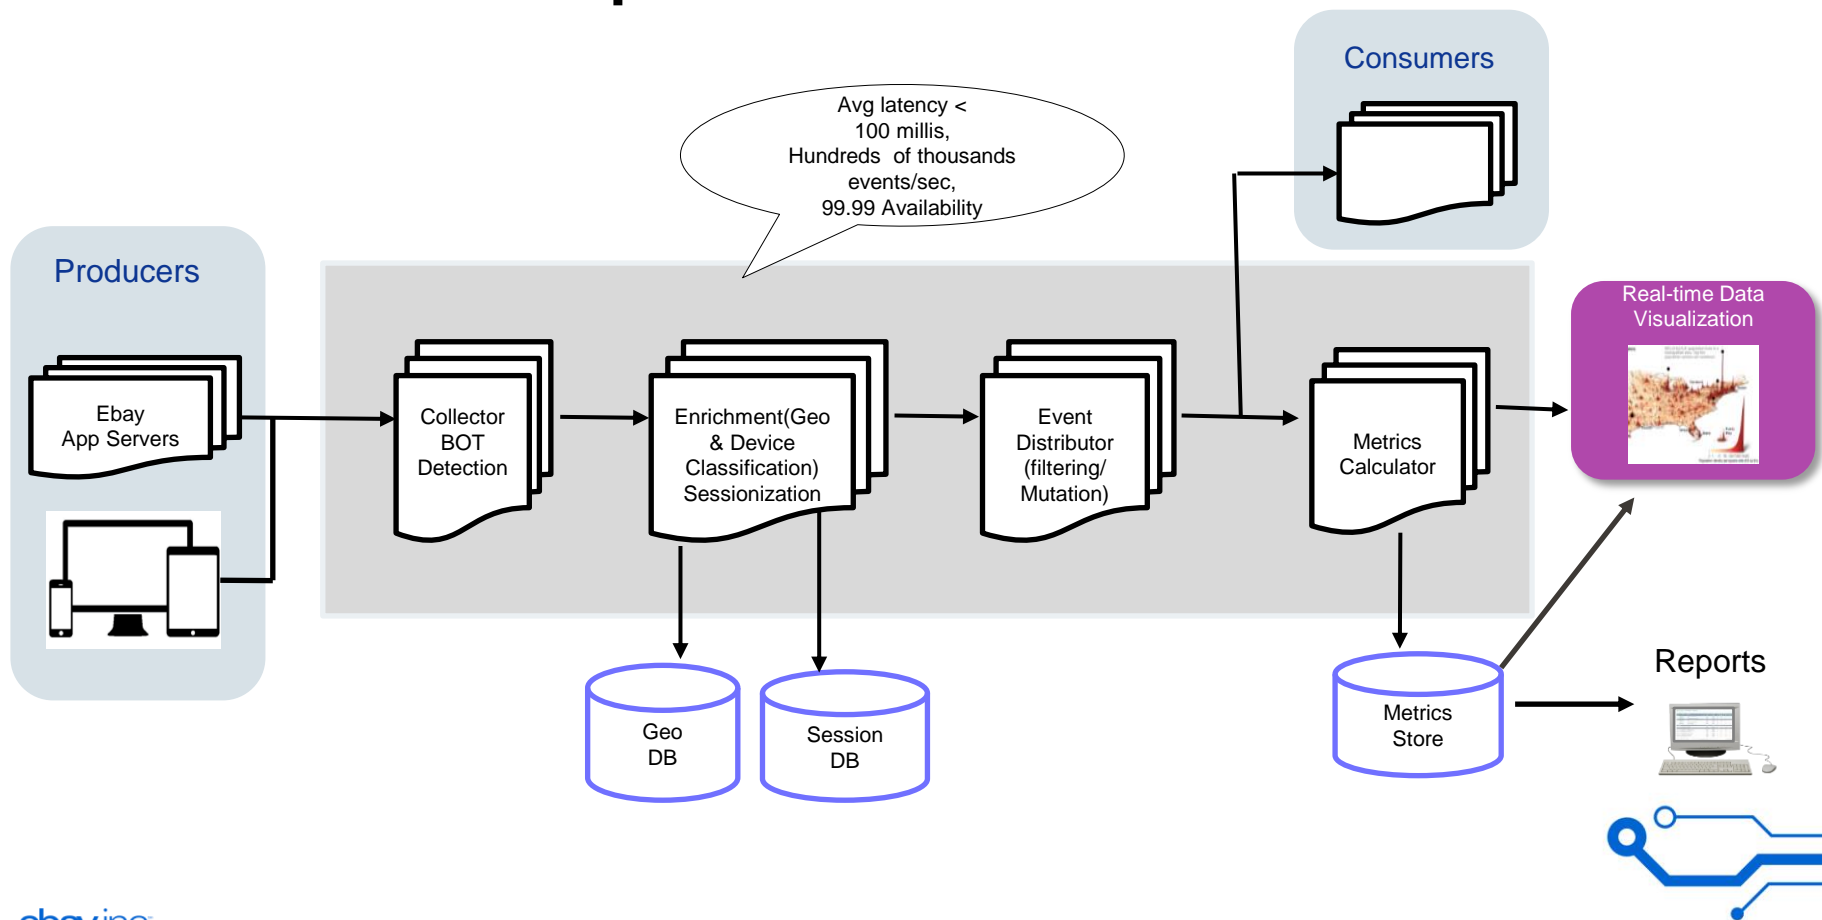


Pulsar Real-time Analytics



- **Complex Event Processing:** SQL on stream data
- **Custom sub-stream creation:** Filtering and Mutation
- **In Memory Aggregation:** Multi Dimensional counting

Pulsar Real Time Pipeline



Key Takeaways

- Creating pipelines declaratively
- SQL driven processing logic with hot deployment of SQL
- Framework for custom SQL extensions
- Dynamic partitioning and flow control
- < 100 millisecond pipeline latency
- 99.99 Availability
- < 0.01% steady state data loss
- Cloud deployable



<http://gopulsar.io>



Q&A
Thanks

